Networking 101

Maël Auzias

ENSIBS - UBS

September 2015

Figure: teaching.auzias.net

Course details

Objectives

- How do *computers* communicate?
- What are the mechanisms **under** an HTTP request or a telegram message?
- Networks are all around us, better study them!

Course details

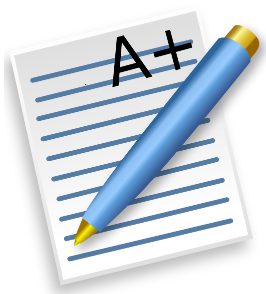

Evaluation

- Short test at the end of each lesson
- Project
- Final exam (1 hour)
- All equal weighting

Material

- Slides available at teaching.auzias.net (github too)

Presentation Outline

Definitions and presentation

- **Network:** an **interconnected** group or system
- **Internet:** world wide **interconnected system of networks** RFC791 (September 1981)
- **IP:** Internet **Protocol** provides the functions necessary to deliver a package of bits from a source to a destination over a network
- **(world wide) Web: network** consisting of a collection of Internet websites using HTTP

Definitions and presentation

- **HTTP:** Hypertext Transfer **Protocol**, application-level protocol for distributed, collaborative, hypermedia information systems draft HTTP2 (July 2014)
- **FTP:** File Transfer **Protocol** promotes sharing of files, encourages the use of remote computers RFC959 (October 1985)
- **RFC:** Request For Comments (Internet Draft (ID), RFC, Internet Standard)

Definitions and presentation

- **Router:** network **hardware** providing routing services
- **Routing: algorithm processed** to decide where to forward a packet
- **Forwarding: action** of moving a packet from one NIC to another
- **NIC:** Network Interface Card
- **Switch (hub):** network **hardware** connecting systems using packet switching
- **Packet switching:** forward-like method regardless of the content (destination-based)
- **NAT:** Network Address Translation, router modifying IP address into another IP address (PAT).

Definitions and presentation

- **Node (network):** any entity that can send packets to/receive packets from a network through a NIC
- **Client: computer** able to send requests to a server
- **Request: application message** destined for a server (*order*)
- **Server: computer** able to respond to a client's requests
- **Response: application message** destined for a client (*result*)
- **Fat client: application** where most functions are processed by the client itself
- **Thin client: application** where most functions are carried out on a central server

Network classification

- **BAN:** Body Area Network
- **PAN:** Personal Area Network
- **(W)LAN:** (Wireless) Local Area Network (home, office, school or airport)
- **MAN:** Metropolitan Area Network, can cover a whole city
- **WAN:** Wide Area Network cover a broad area (Internet)

Topologies

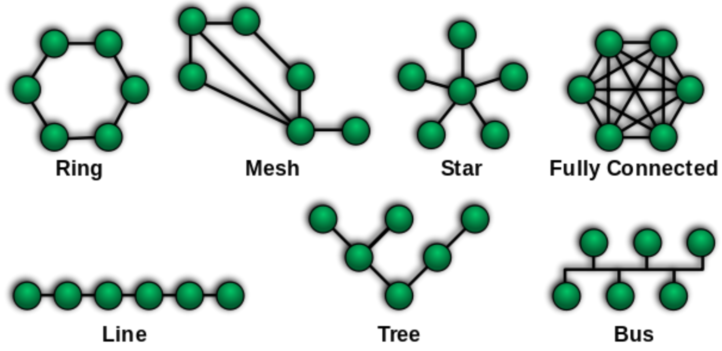

Figure: upload.wikimedia.org

Topologies

- **Point-to-point:** two entities directly connected to each other (tunnel).
- **Ring:** data go around the ring, unidirectional way network.
- **Mesh:** all nodes cooperate in the distribution of data in the network[1].
- **Star:** all messages go through the same central node, reducing network failure.
- **Fully connected:** all nodes are connected to all other nodes.
- **Line:** bidirectional link between two nodes. Node can only send packet going through its neighbors.
- **Bus:** all nodes are connected to the same media. Only one can send a packet at a time, which all others then receive.
- **Tree:** hierarchical topology, such as a binary tree.

[1] Hong Kong protesters used a mesh network to organize (2014)

Bonus

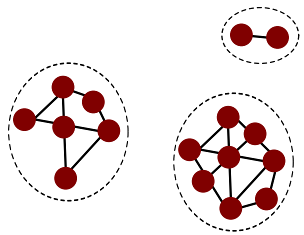

Figure: Disconnected MANET illustration

Bonus

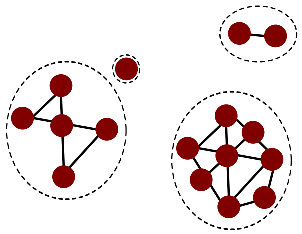

Figure: Store-carry-and-forward

Bonus

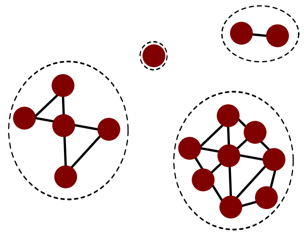

Figure: Store-carry-and-forward

Bonus

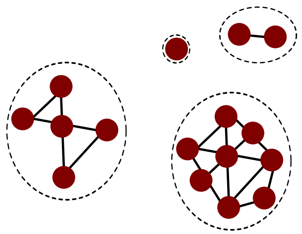

Figure: Store-carry-and-forward

Bonus

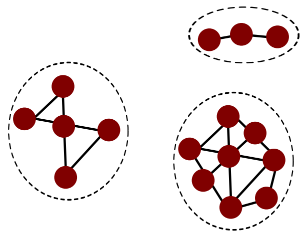

Figure: Store-carry-and-forward

HTTP request/response example

Enter getbootstrap.com in your browser

HTTP request/response example

Enter getbootstrap.com in your browser

Source	Destination	Protocol	Length	Info
192.168.0.48	208.67.222.222	DNS	76	Standard query 0x4797 A getbootstrap.com
208.67.222.222	192.168.0.48	DNS	108	Standard query response 0x4797 A 192.30.252.154 A 192.30.252.153

Figure: DNS request/response

HTTP request/response example

Enter getbootstrap.com in your browser

Source	Destination	Protocol	Length	Info
192.168.0.48	208.67.222.222	DNS	76	Standard query 0x4797 A getbootstrap.com
208.67.222.222	192.168.0.48	DNS	108	Standard query response 0x4797 A 192.30.252.154 A 192.30.252.153

Figure: DNS request/response

Source	Destination	Protocol	Length	Info
127.0.0.1	127.0.0.13	TCP	74	36159 > http [SYN] Seq=0 Win=43690 Len=0 MSS=65495 SACK_PERM=1 TSval=12
127.0.0.13	127.0.0.1	TCP	74	http > 36159 [SYN, ACK] Seq=0 Ack=1 Win=43690 Len=0 MSS=65495 SACK_PERM
127.0.0.1	127.0.0.13	TCP	66	36159 > http [ACK] Seq=1 Ack=1 Win=43776 Len=0 TSval=122257 TSecr=12225
127.0.0.1	127.0.0.13	HTTP	356	GET /index.html HTTP/1.1
127.0.0.13	127.0.0.1	TCP	66	http > 36159 [ACK] Seq=1 Ack=291 Win=44800 Len=0 TSval=122259 TSecr=122
127.0.0.13	127.0.0.1	HTTP	354	HTTP/1.1 200 OK (text/html)
127.0.0.1	127.0.0.13	TCP	66	36159 > http [ACK] Seq=291 Ack=289 Win=44800 Len=0 TSval=122259 TSecr=1
127.0.0.1	127.0.0.13	HTTP	357	GET /favicon.ico HTTP/1.1
127.0.0.13	127.0.0.1	HTTP	565	HTTP/1.1 404 Not Found (text/html)
127.0.0.1	127.0.0.13	TCP	66	36159 > http [ACK] Seq=582 Ack=788 Win=45952 Len=0 TSval=122269 TSecr=1

Figure: HTTP request/response

To read

https://github.com/alex/what-happens-when

- DNS lookup
- ARP process
- Opening of a socket
- TLS handshake
- HTTP protocol
- HTTP Server Request Handle

How do messages reach their destination?

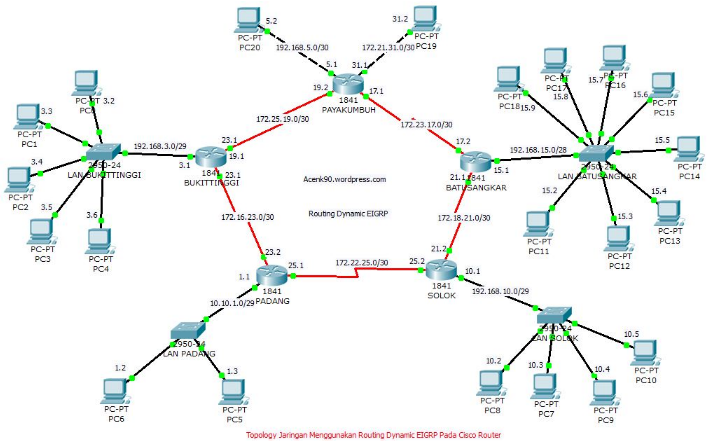

Figure: acenk90.files.wordpress.com

More like this...

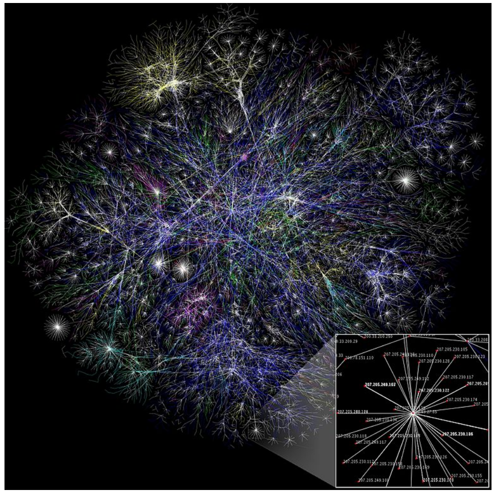

Figure: wikimedia.org

Models overview (OSI and TCP/IP)

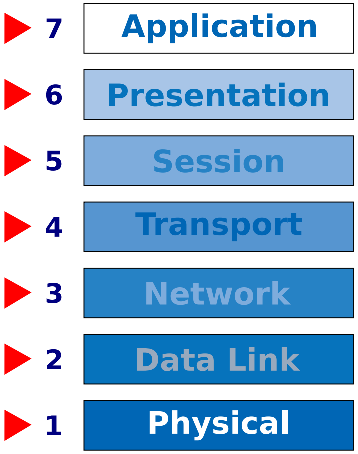

Figure: OSI model

Nth layer communicate with Nth layer..

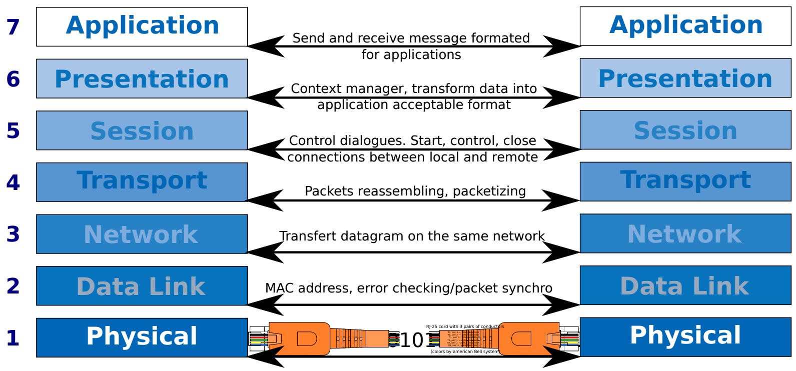

.. thanks to 3-th layers

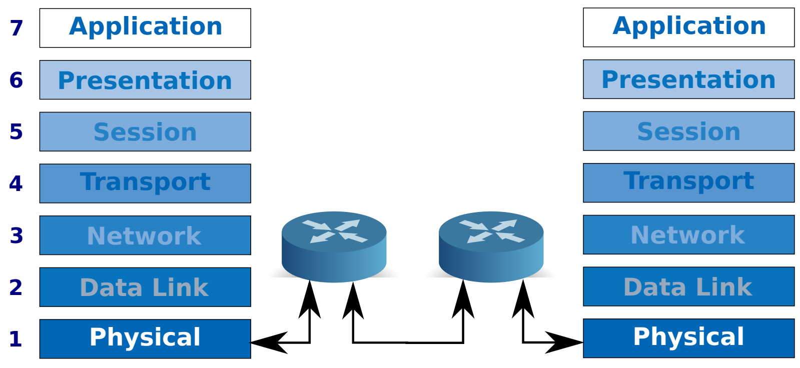

Figure: layers and routing

One single protocol, one single layer

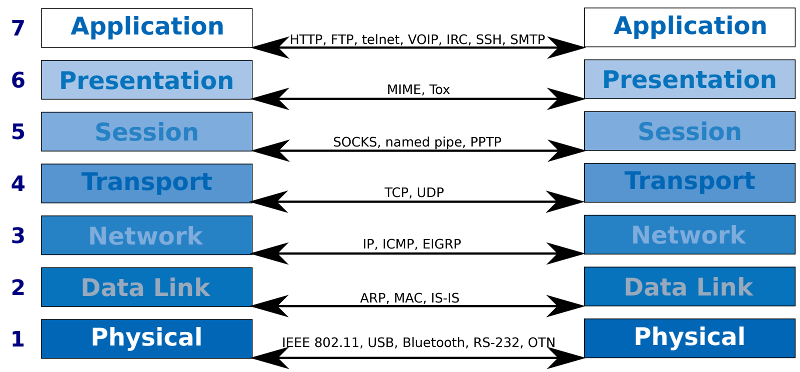

Figure: protocols and layers

Encapsulation

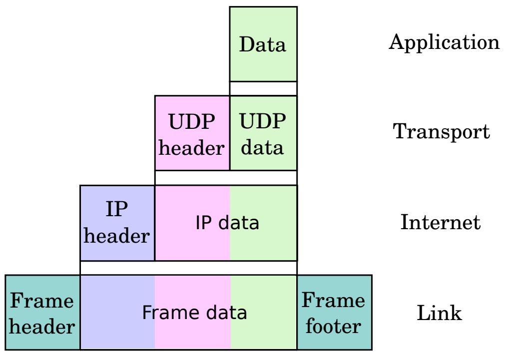

Figure: Encapsulation

Reading

Reading list:

- "Computer Networks" by A Tanenbaum, Andrew S., G ISBN 013162959X
- "Programmation système en C sous Linux" by C Blaess [2] , ISBN 978-2212110548
- http://nmap.org/book/toc.html
- http://blog.nodenexus.com/2014/11/28/a-shark-on-the-network/
- and many many other resources on the Internet freely available[3]! If you can read it, knowledge is reachable[4]!

[2] Translator in French of many man pages
[3] An Introduction to Computer Networks (21: Security) by Peter L Dordal
[4] such as this example of Wireshark using or what-happens-when

Watching

Watching list:

- DEF CON 22 Hacking Conference Presentation By Christopher Soghoian - Blinding The Surveillance State [5]
- any other defcon
- Mr Robot, that's a good serie!

[5] media.defcon.org

Presentation Outline

Aims

- Interface data link layer,
- (De)Encode,
- Transmit: 1 after 0 (after 0 or 1, after 0... or 1)

Hardware medium

- IEEE 802.3 (a.k.a. Ethernet): <100Gbit/s
- IEEE 802.11 (a.k.a. Wi-Fi): <50 Mbit/s (802.11ad goes up to 6.75 Gbit/s)
- IEEE 802.15.1 (a.k.a. Bluetooth): <1 Mbit/s
- IEEE 802.15.4 (a.k.a. ZigBee): <250 kbit/s
- IEEE 802.16 (a.k.a. Wi-Max): <40 Mbit/s
- IEEE 1394 (a.k.a. Firewire): <3200 Mbit/s
- USB, serial port such as RS-232...

Hardware medium: IEEE 802.3 (Ethernet)

Figure: RJ45 connector

Hardware medium: IEEE 802.15.1 (Bluetooth)

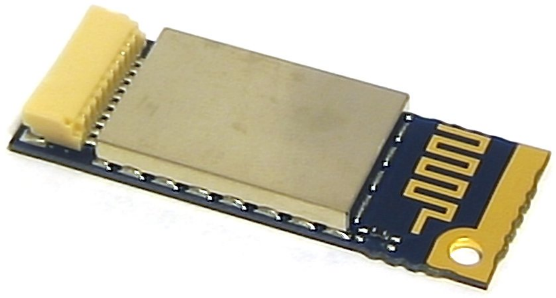

Figure: Bluetooth card

Hardware medium: IEEE 802.15.4 (ZigBee)

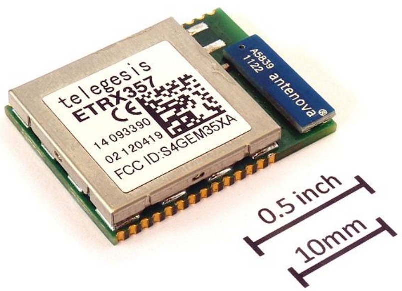

Figure: ZigBee card

Hardware medium: IEEE 802.16 (Wi-Max)

Figure: Wi-Max antenna

Hardware medium: IEEE 1394 (Firewire)

Figure: Firewire connector

Encoding

- **MLT3 (Multi-Level Transmit):** state changes for 1s over 3 levels, stays in the same state for 0s
- **AMI (Alternate Mark Inversion):** state 0 for 0s, state +/-1 for 1s
- **Manchester:** voltage transition (rising/falling edge mean 1/0)
- **BMC (Biphase Mark Code):** change its state for 1s, stay on the same state for 0s
- and so on...

Encoding: Multi-Level Transmit

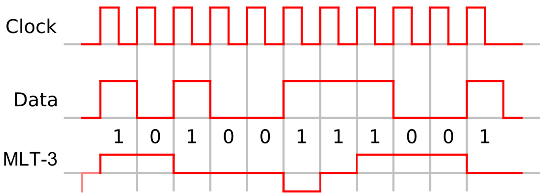

Figure: Multi-Level Transmit

Encoding: Alternate Mark Inversion

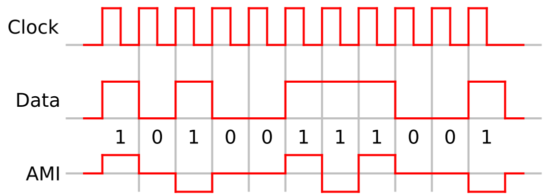

Figure: Alternate Mark Inversion

Encoding: Manchester

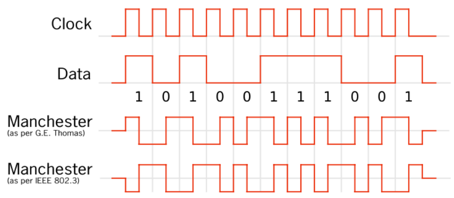

Figure: Manchester

Encoding: Biphase Mark Code

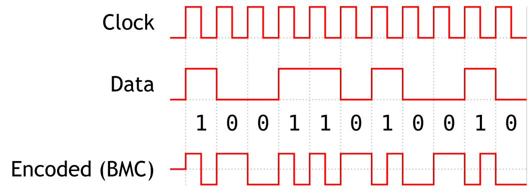

Figure: Biphase Mark Code

Transmitting

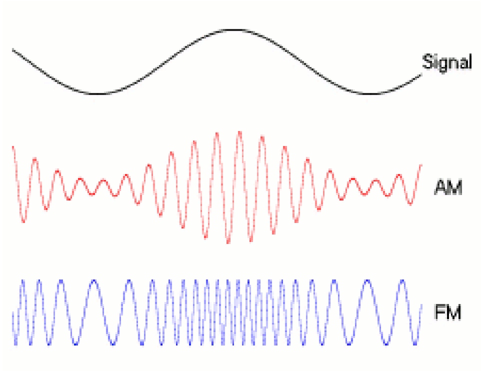

Figure: Amplitude and phase modulation

Error detection

- Repetition (hum...)
- Parity (XOR)
- Checksum
- CRC (Cyclic redundancy check): with a polynomial divison
- Hash
- and so on...

Error correcting

- Repetition (again)
- Hamming
- MDPC (Multidimensional parity-check code)

Correction: MDPC

Raw data to send: 0x01 02 03 04

0x01	0x02	0x03
0x03	0x04	0x07
0x04	0x06	

Figure: Data received with MDPC

Data sent (with MDPC): 0x01 02 03 03 04 07 04 06

Presentation Outline

Aims

- Interface network layer,
- Delivery to unique(?) hardware addresses,
- Framing,
- Data transfer

Layer composition (of its two sublayers)

1. Logical Link Control (LLC):
   - end to end flow control
   - end to end error control
   - (transmitting/receiving) protocols, over MAC sublayer, multiplexing
2. Media Access Control (MAC):
   - physical (hardware) addressing
   - collision detection and retransmission
   - data packet scheduling (and queuing)
   - QoS
   - VLAN

Carrier Sense Multiple Access with Collision Avoidance

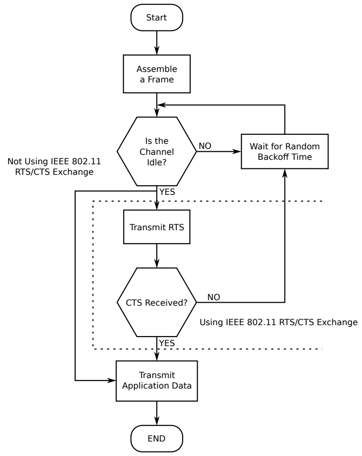

Figure: CSMA CA

Layer 2 Ethernet packet

MAC dest. (6)	MAC src. (6)	VLAN tag* (4)	Ethertype (2)
Payload (42-1500)		Frame check sequence (4)	

Figure: Layer 2 Ethernet packet

optional, Content (size in bytes)

Ethertype 0x	Protocol
0800	IPv4
0806	ARP
0842	Wake-on-LAN
86dd	IPv6

Figure: Data received with MDPC

ARP example

0000	ff	ff	ff	ff	ff	ff	fa	ba	00	ab	ab	af	08	06	00	01
0010	08	00	06	04	00	01	fa	ba	00	ab	ab	af	ac	11	22	37
0020	00	00	00	00	00	00	ac	11	00	f9	00	00	00	00	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00				

Figure: ARP request

MAC address destination MAC address source Ethertype Hardware type Protocol type OpCode (1 request, 2 reply) IP address source IP address destination

ARP example

0000	ff	ff	ff	ff	ff	ff	fa	ba	00	ab	ab	af	08	06	00	01
0010	08	00	06	04	00	01	fa	ba	00	ab	ab	af	ac	11	22	37
0020	00	00	00	00	00	00	ac	11	00	f9	00	00	00	00	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00				

Figure: ARP request

MAC address destination MAC address source Ethertype Hardware type Protocol type OpCode (1 request, 2 reply) IP address source IP address destination

ARP example

0000	fa	ba	00	ab	ab	af	be	be	00	00	eb	eb	08	06	00	01
0010	08	00	06	04	00	01	be	be	00	00	eb	eb	ac	11	00	f9
0020	fa	ba	00	ab	ab	af	ac	11	22	37	00	00	00	00	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00				

Figure: ARP reply

MAC address destination MAC address source Ethertype Hardware type Protocol type OpCode (1 request, 2 reply) IP address source IP address destination

ARP example

0000	fa	ba	00	ab	ab	af	be	be	00	00	eb	eb	08	06	00	01
0010	08	00	06	04	00	01	be	be	00	00	eb	eb	ac	11	00	f9
0020	fa	ba	00	ab	ab	af	ac	11	22	37	00	00	00	00	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00				

Figure: ARP reply

MAC address destination MAC address source Ethertype Hardware type Protocol type OpCode (1 request, 2 reply) IP address source IP address destination

Presentation Outline

Aims

- Interface transport layer,
- Host addressing,
- End-to-end packet transmission (data link? Connectionless? Switch? Router?),
- Routing, load balancing

Concepts

- IP addressing fundamentals,
- Classfull IP addressing,
- Subnet and VLSM (Variable length subnet masks),
- CIDR (Classless inter-domain routing),
- Routing,
- IPv6.

IP addressing fundamentals

IP address

32 bits (4x4 bytes)

mask	
Networks part	Host part

Figure: IP address parts

IP addressing fundamentals

Masks

- Separates network and host bits,
- MSB are **always** ones and then zeros! 255.254.255.0 is not possible,
- Indicates how many bits are used for the network part:
  - A 8-bit mask leaves 24 bits for the hosts,
  - A 16-bit mask leaves 16 bits for the hosts,
  - A 24-bit mask leaves 8 bits for the hosts,
  - A N-bit mask leaves 32-N bits for the hosts.
- Two different masks (differences seen further on):
  - Network mask,
  - Subnet mask.

IP addressing fundamentals

IP address

32 bits (4x4 bytes)	
Networks part	Host part

Figure: IP address parts and mask

IP addressing fundamentals

IP address

32 bits (4x4 bytes)	
ones mask	zeros mask
Networks part	Host part

Figure: IP address parts and mask

IP addressing fundamentals

Is that an address?

- Network address,
- Hosts,
- Broadcast address.

Within the same network

- All addresses have the same network bits,
- Network address has zeros for host bits: x.x.x.0*,
- All hosts have different host bits: x.x.x.[0-1]*,
- Broadcast address has ones for host bits: x.x.x.1*.

IP addressing fundamentals

Mask /24 254 hosts	255 11111111	255 11111111	255 11111111	0 00000000
Network address	192 11000000	168 10101000	1 00000001	0 00000000
First host	192 11000000	168 10101000	1 00000001	1 00000001
Last host	192 11000000	168 10101000	1 00000001	254 11111110
Broadcast address	192 11000000	168 10101000	1 00000001	255 11111111

Figure: IP address example 1

IP addressing fundamentals

Mask /16 65.534 hosts	255 11111111	255 11111111	0 00000000	0 00000000
Network address	172 10101100	64 01000000	0 00000000	0 00000000
First host	172 10101100	64 01000000	0 00000000	1 00000001
Last host	172 10101100	64 01000000	255 11111111	254 11111110
Broadcast address	172 10101100	64 01000000	255 11111111	255 11111111

Figure: IP address example 2

IP addressing fundamentals

Formula: how many hosts with an N-bit mask?

$2^{32-N} - 2$, the -2 moves out network and broadcast addresses which are not hosts.

- 24-bit mask: $2^{32-24} - 2 = 2^8 - 2 = 254$ hosts
- 16-bit mask: $2^{32-16} - 2 = 2^{16} - 2 = 65.534$ hosts
- 8-bit mask: $2^{32-8} - 2 = 2^{24} - 2 = 16.777.214$ hosts

IP addressing fundamentals

Public addresses

- Most IP addresses
- Registered ISP and large organizations inherit blocks of public addresses from IANA[6]
- Usage of unregistered public addresses is forbidden.

Private addresses

- Privates addresses are A, B and C classes (not all, see after)
- No registration needed
- Not routed across the Internet
- Proxy, NAT and private addresses solved IPv4 shortage.

[6] Internet Assigned Numbers Authority

Classful IP Addressing

Class	A	B	C
First octet	1 - 126	128 - 191	192 - 223
First octet 0b	0*	10*	110*
Network mask	255.0.0.0 /8	255.255.0.0 /16	255.255.255.0 /24
IP addresses range	1.0.0.0 126.0.0.0	128.0.0.0 191.255.0.0	192.0.0.0 223.255.255.0
Private range	10.0.0.0 10.255.255.255	172.16.0.0 172.31.255.255	192.168.0.0 192.168.255.0
Number of hosts	16.777.214	65.534	254

Figure: Three main classes

Where did 127.0.0.0/8 go ?!

Classful IP Addressing

Class D

- First octet: 224 - 239
- First octet pattern: 1110*
- These IP addresses are multicast addresses.

Class E

- Everything left
- Experimental class.

Classful IP Addressing

Reserved addresses

- 0.0.0.0 used in routing (seen further)
- 127.0.0.0/8: loopback addresses (127.0.0.1 - 127.255.255.254).

Classful IP Addressing

- Class A (16 m-addresses) and B (65 k-adresses) are too large!
- Class C (254 addresses) is manageable. A and B are not, and then not fully utilized... That's a waste of IP addresses!

Three means to limit the number of nodes on a network (regardless of the class) and, thus, improve manageability:

- Subnet,
- VLSM (Variable Length Subnet Mask),
- CIDR (Classless Inter-Domain Routing).

Subnet and VLSM

- Class A (16 m-addresses) and B (65 k-adresses) are too large!
- Class C (254 addresses) is manageable. A and B are not, and then not fully utilized... That's a waste of IP addresses!

Subnet and VLSM

Mask /16 65.534 hosts	255 11111111	255 11111111	0 00000000	0 00000000
Network address	172 10101100	64 01000000	0 00000000	0 00000000
First host	172 10101100	64 01000000	0 00000000	1 00000001
Last host	172 10101100	64 01000000	255 11111111	254 11111110
Broadcast address	172 10101100	64 01000000	255 11111111	255 11111111

Figure: IP address example 2

Subnet and VLSM

Mask /12 1.048.574 hosts	255 11111111	240 11110000	0 00000000	0 00000000
Network address	172 10101100	64 01000000	0 00000000	0 00000000
First host	172 10101100	64 01000000	0 00000000	1 00000001
Last host	172 10101100	79 01001111	255 11111111	254 11111110
Broadcast address	172 10101100	79 01001111	255 11111111	255 11111111

Figure: IP address example 3

Subnet and VLSM

Mask /10 4.194.302 hosts	255 11111111	192 11000000	0 00000000	0 00000000
Network address	172 10101100	64 01000000	0 00000000	0 00000000
First host	172 10101100	64 01000000	0 00000000	1 00000001
Last host	172 10101100	127 01111111	255 11111111	254 11111110
Broadcast address	172 10101100	127 01111111	255 11111111	255 11111111

Figure: IP address example 4

Subnet and VLSM

Mask /31 0 host	255 11111111	255 11111111	255 11111111	254 11111110
Network address	172 10101100	64 01000000	0 00000000	254 11111110
First host	172 10101100	64 01000000	0 00000000	? 1111111?
Last host	172 10101100	64 01000000	255 00000000	? 1111111?
Broadcast address	172 10101100	64 01000000	255 00000000	255 11111111

Figure: IP address example 5

Subnet and VLSM

Mask /30 2 hosts	255 11111111	255 11111111	255 11111111	252 11111100
Network address	172 10101100	64 01000000	0 00000000	252 11111100
First host	172 10101100	64 01000000	0 00000000	253 11111101
Last host	172 10101100	64 01000000	255 00000000	254 11111110
Broadcast address	172 10101100	64 01000000	255 00000000	255 11111111

Figure: IP address example 6

Netmask		CIDR	hosts
255.255.255.255	11111111.11111111.11111111.11111111	/32	Unusable
255.255.255.254	11111111.11111111.11111111.11111110	/31	Unusable
255.255.255.252	11111111.11111111.11111111.11111100	/30	2
255.255.255.248	11111111.11111111.11111111.11111000	/29	6
255.255.255.240	11111111.11111111.11111111.11110000	/28	14
255.255.255.224	11111111.11111111.11111111.11100000	/27	30
255.255.255.192	11111111.11111111.11111111.11000000	/26	62
255.255.255.128	11111111.11111111.11111111.10000000	/25	126
255.255.255.0	11111111.11111111.11111111.00000000	/24	254
255.255.254.0	11111111.11111111.11111110.00000000	/23	510
255.255.252.0	11111111.11111111.11111100.00000000	/22	1.022
255.255.248.0	11111111.11111111.11111000.00000000	/21	2.046
255.255.240.0	11111111.11111111.11110000.00000000	/20	4.094
255.255.224.0	11111111.11111111.11100000.00000000	/19	8.190
255.255.192.0	11111111.11111111.11000000.00000000	/18	16.382
255.255.128.0	11111111.11111111.10000000.00000000	/17	32.766
255.255.0.0	11111111.11111111.00000000.00000000	/16	65.534
255.254.0.0	11111111.11111110.00000000.00000000	/15	131.070
255.252.0.0	11111111.11111100.00000000.00000000	/14	262.142
255.248.0.0	11111111.11111000.00000000.00000000	/13	524.286
255.240.0.0	11111111.11110000.00000000.00000000	/12	1.048.574
255.224.0.0	11111111.11100000.00000000.00000000	/11	2.097.152
255.192.0.0	11111111.11000000.00000000.00000000	/10	4.194.302
255.128.0.0	11111111.10000000.00000000.00000000	/9	8.388.606
255.0.0.0	11111111.00000000.00000000.00000000	/8	16.777.214
254.0.0.0	11111110.00000000.00000000.00000000	/7	33.554.430
252.0.0.0	11111100.00000000.00000000.00000000	/6	67.108.862
248.0.0.0	11111000.00000000.00000000.00000000	/5	134.217.726
240.0.0.0	11110000.00000000.00000000.00000000	/4	268.435.454
224.0.0.0	11100000.00000000.00000000.00000000	/3	536.870.910
192.0.0.0	11000000.00000000.00000000.00000000	/2	1.073.741.822
128.0.0.0	10000000.00000000.00000000.00000000	/1	2.147.483.646
0.0.0.0	00000000.00000000.00000000.00000000	/0	IP space

Figure: Subnet mask cheat sheet

CIDR

Classless Inter-domain Routing?

- Wait! What is routing?

Routing Principles

Algorithms are processed to decide where to forward a packet

Any router must

- know where any packet should be directed
- send directly the packets to the destination if the router and the destination are on the same (sub)network

Any node

- on any network can communicate directly with all the nodes within the same network
- can connect to any node using its gateway
- needs to be aware of its gateway to communicate with nodes on other networks

Routing Principles

Route

- Destination
- Gateway (next hop)
- Masks
- Metric
- Interface

```
>sudo route -n
Kernel IP routing table
Destination     Gateway         Genmask         Flags Metric Ref    Use Iface
0.0.0.0         192.168.0.254   0.0.0.0         UG    0      0        0 eth0
192.168.0.0     0.0.0.0         255.255.255.0   U     0      0        0 eth0
```

Figure: Routing table

Routing Principles

```
>sudo route -n
Kernel IP routing table
Destination     Gateway         Genmask         Flags Metric Ref    Use Iface
0.0.0.0         192.168.0.254   0.0.0.0         UG    0      0        0 eth0
192.168.0.0     0.0.0.0         255.255.255.0   U     0      0        0 eth0
```

Figure: Routing table

0.0.0.0 ?

- Default destination
- Default (sub)network(s)
- Default route
- Default gateway

Routing Principles

Example

what would the routing table of this router look like?

Routing Principles

Static or dynamic ?

We will see this later

CIDR

Combine 2+ networks' into one bigger to ease routing.

Classless Inter-domain Routing?

- Can a routing table having both (192.168.0.0/24, E0), (192.168.1.0/24, E0), (10.0.0.0/8, S0) be shorten?
- Can a routing table having both (192.168.0.0/24, E0), (192.168.1.0/24, E0), (192.168.8.0/24, E0), (10.0.0.0/8, S0) be shorten?
- Can a routing table having both (192.168.0.0/24, E0), (192.168.4.0/24, E0), (192.168.1.0/24, E1), (10.0.0.0/8, S0) be shorten?

Routing Protocol

- RIP: Routing Information Protocol
- OSPF: Open Shortest Path First
- EIGRP: Enhanced Interior Gateway Routing Protocol

Routing Protocol

RIP v1

- Classful routing
- Periodic updates (30 sec) ..
- ..by broadcasting (!)
- Metric is hop-count (max = 15, infinite = 16)
- Timer (180 sec) to tag route as invalid (metric = 16)
- no subnet, no VLSM, no CIDR, no router authentication

Routing Protocol

RIP v2

- Classless routing
- Multicast (224.0.0.9)
- VLSM support
- Route summarization
- "Authentication" (MD5)

RIPng is the next RIP version for support of IPv6

Routing Protocol

1. Router coming online broadcasts Request message
2. RIP Routers send **broadcasts** Response messages with their routing table
3. When Update timers (from other routers) expire, its routing table[7] is sent again
4. When Invalid timer expires, the metric of the route is set to 16 (unreachable)
5. When Flush timer expires, the 16-metric routes are removed from the routing table
6. When a new router (or new metric) is sent, a Hold-down timer is started to stabilize the network.

[7] not always the whole table

Routing Protocol

OSPF

- Classless
- IPv4 and IPv6
- VSLM
- CIDR
- Build a topology of the network
- Dijkstra
- Metric = f(hop-count, bandwidth, link reliability)
- Subdivided into area (a 32-bit number)
- Multicast
- Authentication support (update only from trusted routers)

Routing Protocol

EIGRP

- Enhanced IGRP (to support classless routing)
- IPv4 and IPv6
- VSLM
- CIDR
- Build a topology of the network
- Dijkstra
- Metric = f(bandwidth, load, delay, reliability)
- Authentication support

IPv6 - Aims

- Support billions of hosts (even with inefficient IP addressing)
- Reduce routing table size
- Simplified protocol to allow routers to process packets faster
- Better security
- Better real-time QoS
- Better multicast diffusion (scope)
- Able to move without changing IP address
- Give the protocol the ability to evolve
- Give the protocol the ability to coexist with newer version

IPv4 vs IPv6

- not compatible
- IPv4 address: 4 octets, IPv6: 16 octets ($2^{128} = 3x10^{138}$)
- Packet Header, IPv6: 7 fields, IPv4:13 (faster to process)
- IP options: some required options are now optional (faster to process)
- Notation:
  - 8000:0000:0000:0000:0123:4567:89AB:CDEF
  - 8000::0123:4567:89AB:CDEF
  - ::192.168.2.3
- Unicast address format:

bits	48 (or more)	16 (or fewer)	64
field	routing prefix	subnet id	interface identifier

Figure: Unicast IPv6 address format

IPv4 vs IPv6

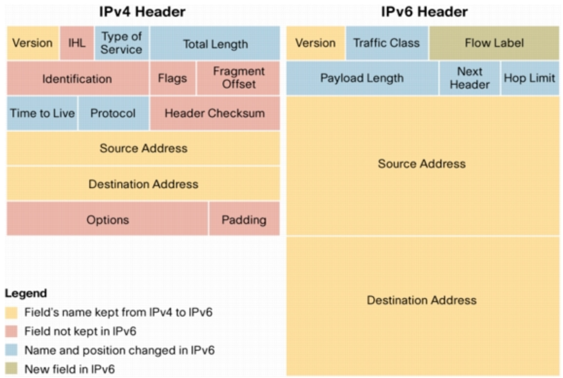

Figure: IPv4 and IPv6 headers (www.cisco.com)

IPv6 - Header

- **Version (4 bits):** 0b0110
- **Traffic class (8 bits):** 6-MSB for differentiated services[8], 2-LSB for ECN[9]
- **Flow label (20 bits):** routers are supposed to use the same path for the same flow (thus, destination do not need to re-order packets)
- **Payload length (16 bits):** packet length minus its header length

[8] multimedia or http
[9] Explicit Congestion Notification (RFC 3168)

IPv6 - Header

- **Next header (8 bits):** specifies the transport layer protocol, also indicates (if any) extension header that follows.
- **Hop limit (8 bits):** Hop count (discussion was to use a duration instead, but router implementations would be much more complex)

Optional IPv6 headers offer the possibility to

- specify the route of the datagram
- include authentication data
- include fragmentation parameters
- and so on...

IPv6 - Anecdotes

- IPv6 address length could have been 8 bytes, or 20 bytes, or even variable
- Hop count max value (255) is considered, by some, not enough
- Removing IPv4 checksum is *as safe as removing brakes from a car*
- Different national laws on encryption disallow a real secure transport layer

IPv6 - Adoption

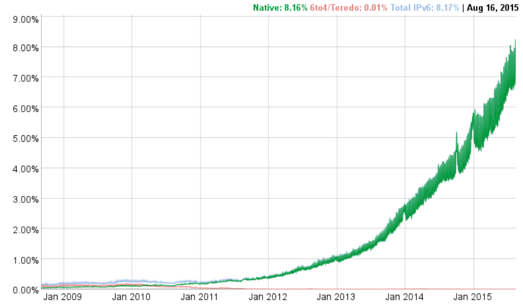

Figure: IPv6 adoption (among Google users)[10]

- **2014** Belgium: 28%, USA and Germany: 11%
- **2015** Belgium: 36%, USA: 21% and Germany: 18%

[10] https://www.google.com/intl/en/ipv6/statistics.html

Presentation Outline

Aims

- Interface session layer,
- Reliable end-to-end communication,
- Order and reassemble received packets (if needed),
- Flow control,
- Congestion avoidance (if supported by protocol),
- Multiplexing

Application identification

Socket address

- Node identification is made by IP address,
- Application identification is made by node identification...
- ... and a port. Number between 0 and 65535. (1-1024: root privilege)
  - ip.ad.dr.ess:port

Port	Protocol
21	FTP
22	SSH
23	Telnet
25	SMTP
80	HTTP
443	HTTPS
465	SMTPS
631	IPP
1194	OpenVPN
3128, 8080	Web Proxy
9418	git
23399	Skype

Figure: Default port for well known protocol

TCP header

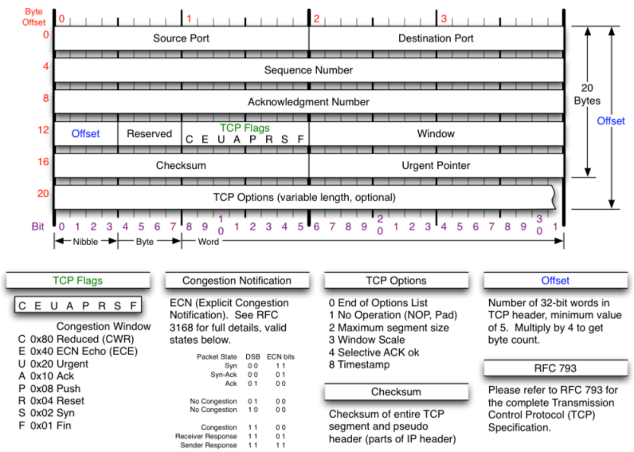

Figure: nmap.org: TCP header

UDP header

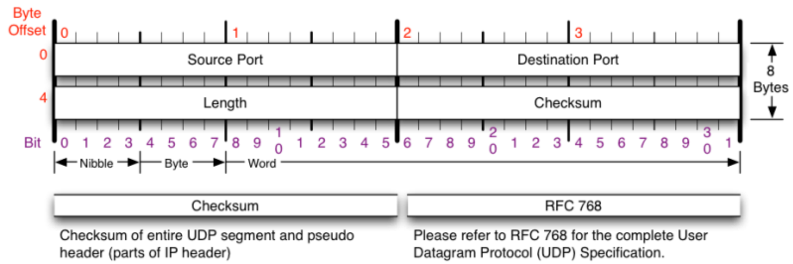

Figure: nmap.org: UDP header

Socket Primitives (TCP)

Order	Primitive	Meaning
1	SOCKET	Creates a new communication endpoint
2	BIND	Links local IP address to the socket (for server)
3	LISTEN	Signs up for incoming connections
4	ACCEPT	Blocking call till a connection attempt occurs
-	CONNECT	**Tries** to connect to another communication endpoint
-	SEND	Sends data through the established connection
-	RECEIVE	Receives data through the established connection
last	CLOSE	Releases the connection (do not mistake shutdown and close.)

Figure: TCP primitives

A socket does not have an IP address until it is bound, just an allocation in the transport entity. A server must listen before any client can connect.
close() a socket does not send the closing stream three handshake, *shutdown*() does. *fork*() is needed, *poll*() and *select*() can be used too.

What are these ?

- **Frame**: Physical layer representation
- **Datagram**: UDP[11] or IP packet (IP datagram, UDP datagram)
- **Segment**: TCP data unit
- **PDU**: Protocol Data Unit, generic term.
- **Fragment**: Any data unit **fragmented**

[11] User **Datagram** Protocol

I hope you liked it and learnt something new !

Figure: teaching.auzias.net